

# UC Davis STA 208 2016 Spring Midterm Exam

Wenhao Wu, 998587583

May 1, 2016

## 1 Data Description

This dataset contains  $N = 3089$  observations of  $p = 4$  features. Among this observations  $N_0 = 1089$  observations are labeled with the response  $y = 0$  and  $N_1 = 2000$  observations with  $y = 1$ , and the formers are located in the second half of the data file while the latter in the first half. The scatterplot matrix of the dataset is shown in Fig. 1

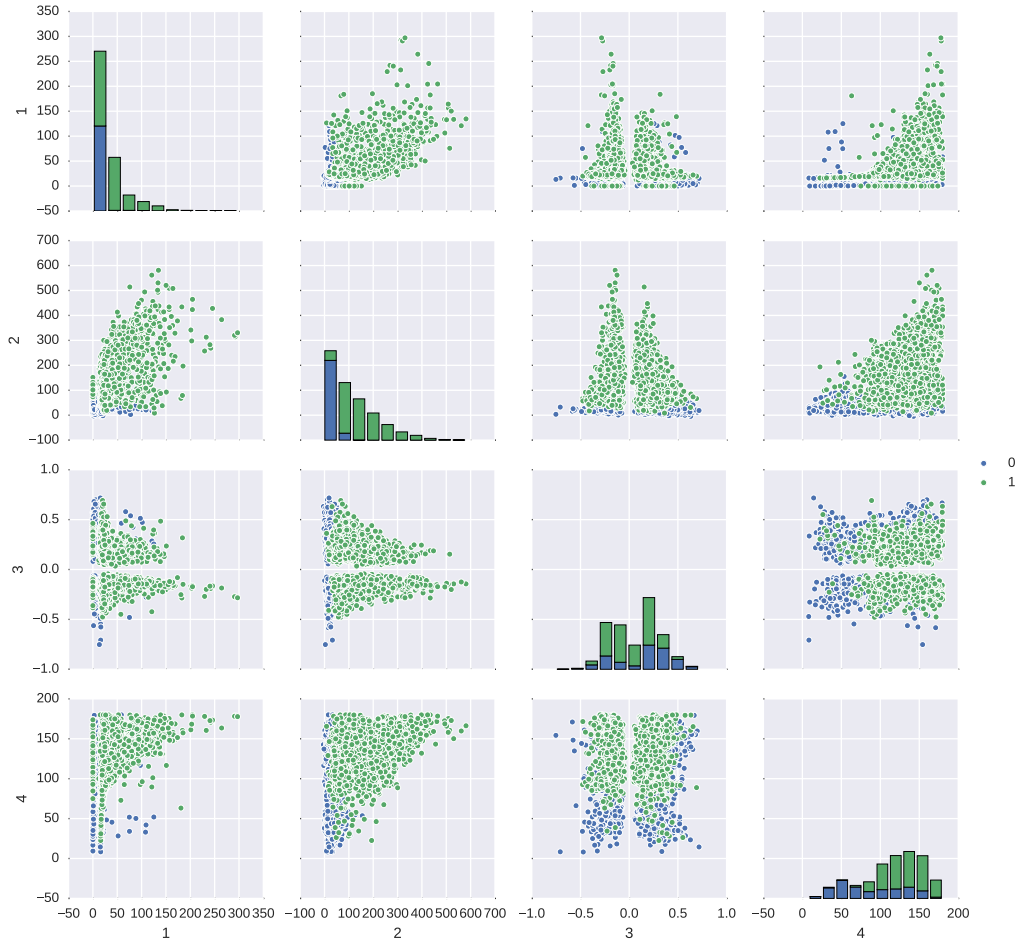


Figure 1: The scatterplot matrix of the dataset.

## 2 Algorithm Design

In order to train and test our classifier, firstly we set aside %25 of the observation as a testing set. The parameters of our classifiers are tuned using 10-fold cross-validation using the rest %75 observations. Prior to the training, the data is preprocessed so that each feature has zero mean and unit variance.

Since the dimension of the input variables is rather low and the number of observations is relatively small, we first try to tune a simple  $k$ -Nearest-Neighbor (KNN) classifier on the 4 given features. The results are shown in Fig. 2. The 10-fold cross validation suggests that when  $k \in [3, 30]$  the validation score is maximized, where the KNN classifier results in a decent test score of  $\sim 0.95$ . We select  $k = 15$  for our KNN classifier.

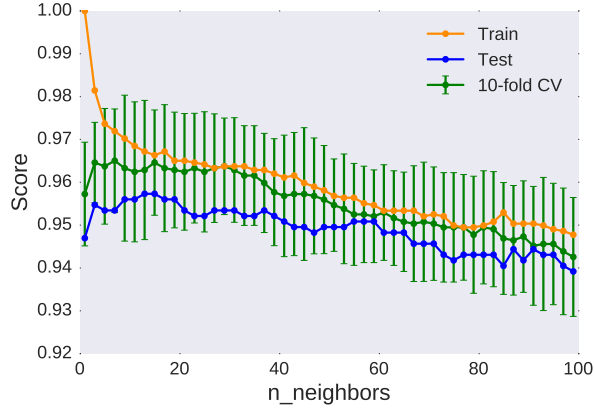


Figure 2: Tuning the KNN classifier.

Although KNN results in a decent performance accuracy, we would like to try other machine learning methods with better interpretability. Here we tried decision tree model in which the depth of the tree is tuned with cross-validation. The results is shown in Fig. 3. It appears that the cross-validation favors a depth of 3 which also results in a test score of  $\sim 0.95$ , for which the decision tree fitted on the training set is shown in Fig. 4.

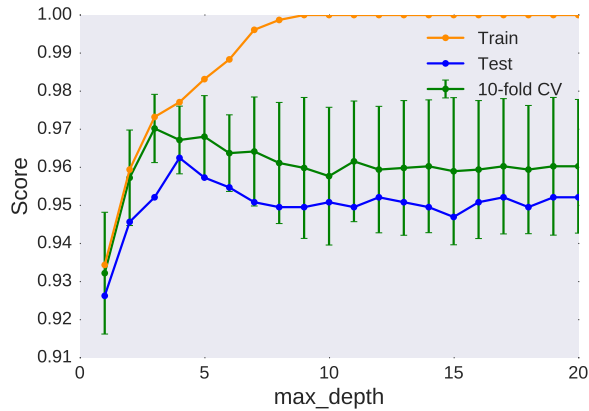


Figure 3: Tuning the decision tree classifier.

Next we try some more rigorous (mathematically) machine learning techniques. Firstly a discriminative model. Specifically, we applied the logistic-regression method where the regularization parameter  $C$  is chosen with CV. The results are shown in Fig. 5. It appears that logistic regression does not work as good as other classifiers with a test score of 0.944 for  $C = 1$ .

We would also like to see whether the performance can be improved with kernel techniques and that promotes sparsities. We try to tune a SVM classifier with rbf kernel, in which both the regularization parameter  $C$  and the kernel parameter  $\gamma$  are chosen with CV. The results are shown in Fig. 6. The SVC achieve a tiny-slightly

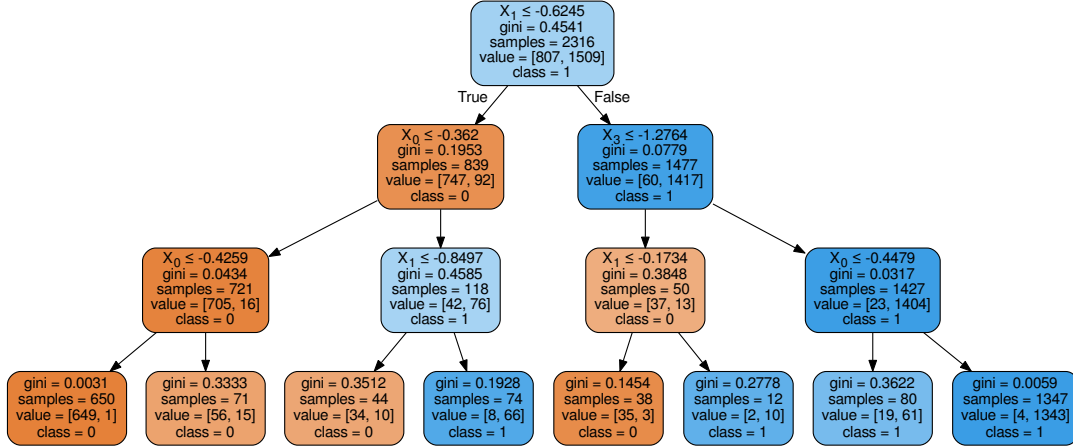


Figure 4: The trained decision tree classifier.

Table 1: Test performances of different classifiers.

Classifier	false positive	false negative	precision	reall	score
KNN	0.06383	0.03055	0.9382	0.9695	0.9573
Decision Tree	0.08511	0.02648	0.9196	0.9735	0.9521
Log-Reg	0.06738	0.04888	0.9338	0.9511	0.9444
SVM	0.05674	0.03055	0.9447	0.9695	0.9599
Random Forest	0.05674	0.02444	0.9450	0.9756	0.9638
RVM	0.06738	0.02851	0.9351	0.9715	0.9573
Variational Log-Reg	0.06738	0.04888	0.9338	0.9511	0.9444

better test score of 0.96 which should not be considered as a significant performance gain. We select  $C = 1$  and  $\gamma = 1$  for our SVM classifier.

The classifiers we have tried so far are tuned with cross-validation, and as we can see the choice of a single best classifier is usually unclear. To reduce the variance of our classifier, we also tried the ensemble learning technique of random forest. We fix the `max_depth=3` and trained a random forest of 10 trees. The resulting test score is 0.968 which turns out to be a little higher than the other classifiers.

As opposed to the frequentist approaches mentioned above, an arguably more insightful and principled way of regularization is the Bayesian approach. In this study, we also tried two Bayesian classifiers, namely the relevance vector machine (RVM) and the variational logistic regression (VLR) implemented by Amazasp Shaumyan. For the RVM classifier, we adopt the rbf kernel with  $\gamma = 1$ . The RVM and VLR achieve test scores of 0.961 and 0.944, respectively, which is almost the same as their non-Bayesian counterparts.

### 3 Numerical Results and Conclusion

The performance of all the classifiers mentioned in the last section are summarized in Table 1. We decide to use the random forest of 10 trees with max depth of 3 as our final classifier. A score of around 0.96 is expected from this classifier.

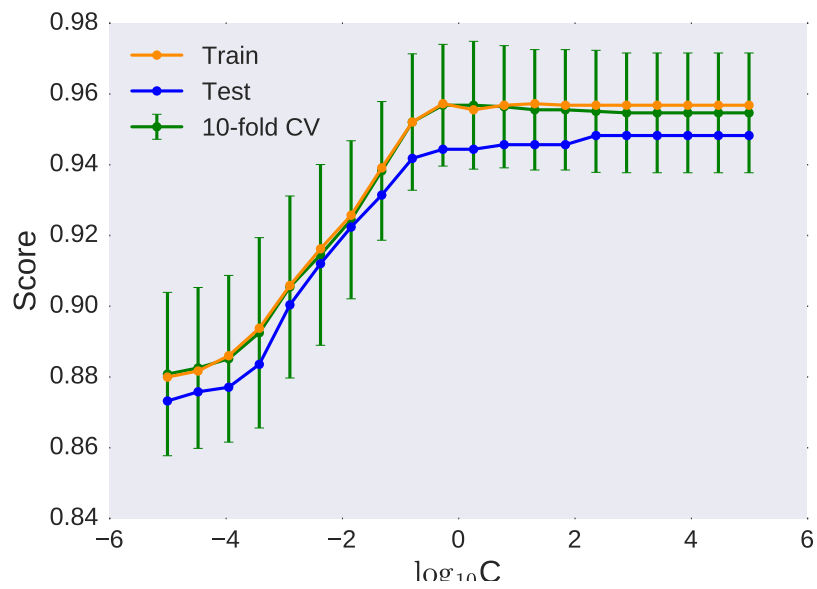


Figure 5: Tuning the logistic regression classifier.




Figure 6: Tuning the SVM classifier.