

## **STA 208: Assignment #2**

Due on Monday, April 11, 2016

*Prof. James Sharpnack MW 12:00 - 2:00 P.M.*

**Wenhao Wu**

## Contents

<b>1.</b>	<b>3</b>
<b>2.</b>	<b>4</b>
<b>3.</b>	<b>5</b>

**1.**

The following losses are used as surrogate losses for large margin classification. Demonstrate if they are convex or not, and follow the instructions.

(a) exponential loss:  $\phi(x) = e^{-x}$

**Answer:** Since  $d^2\phi(x)/dx^2 = e^{-x} > 0$ , this function is convex.

(b) truncated quadratic loss:  $\phi(x) = (\max\{1 - x, 0\})^2$

**Answer:** We can rewrite  $\phi(x) = (\max\{1 - x, 0\})^2$

$$\phi(x) = \begin{cases} (1 - x)^2 & x < 1, \\ 0 & \text{otherwise.} \end{cases} \quad (1)$$

Then the convexity of  $\phi(x)$  can be proved by the definition

- $\forall x_1 < 1, x_2 < 1$  or  $\forall x_1 \geq 1, x_2 \geq 1$ , since both  $f(x) = (1 - x)^2$  and  $g(x) = 0$  are convex functions, we have

$$\phi(tx_1 + (1 - t)x_2) \leq t\phi(x_1) + (1 - t)\phi(x_2) \quad (2)$$

$$\forall t \in [0, 1].$$

- $\forall x_1 < 1, x_2 \geq 1$ , since  $\phi(x)$  is monotonically decreasing,

$$\phi(tx_1 + (1 - t)x_2) \leq \phi(tx_1) = t\phi(x_1) + (1 - t)\phi(x_2) \quad (3)$$

Therefore  $\phi(x)$  is convex.

(c) hinge loss:  $\phi(x) = \max\{1 - x, 0\}$

**Answer:** Since both  $f(x) = 1 - x$  and  $g(x) = 0$  are convex functions,  $(\max\{1 - x, 0\})$ ,  $\phi(x)$  is also convex.

(d) sigmoid loss:  $\phi(x) = 1 - \tanh(\kappa x)$ , for fixed  $\kappa > 0$

**Answer:** Since

$$\frac{d\phi}{dx} = \frac{-4\kappa e^{2\kappa x}}{(1 + e^{2\kappa x})} = \kappa\phi(x)(\phi(x) - 2) \quad (4a)$$

$$\frac{d^2\phi}{dx^2} = 2\kappa^2\phi(x)(\phi(x) - 2)(\phi(x) - 1) \quad (4b)$$

when  $x < 0$ ,  $d^2\phi/dx^2 < 0$ , therefore  $\phi(x)$  is not convex.

(e) Plot these as a function of  $x$ .

**Answer:**

```
import numpy as np
import matplotlib.pyplot as plt
import matplotlib as mpl
%matplotlib qt

x = np.linspace(-2, 2, num=50)
kappa = 1 # The parameter for the sigmoid loss
```

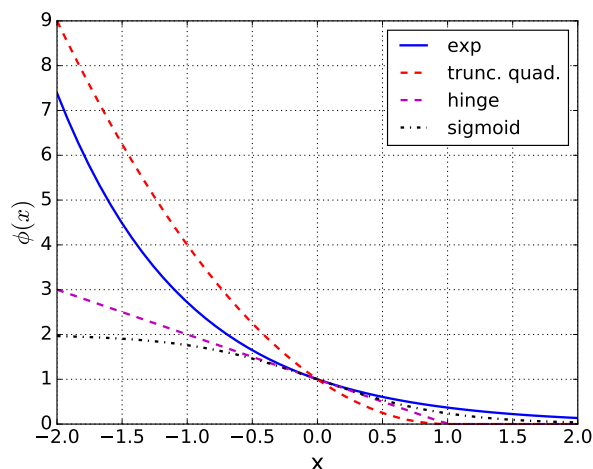


Figure 1: Comparison of different loss functions.

```
axis_font = {'size': '20'}
mpl.rcParams['xtick.labelsize'] = 16
mpl.rcParams['ytick.labelsize'] = 16

plt.plot(x, np.exp(-x), 'b-', linewidth=2, label="exp")
plt.plot(x, (1-x).clip(0) ** 2, 'r--', linewidth=2, label="trunc. quad.")
plt.plot(x, (1-x).clip(0), 'm--', linewidth=2, label="hinge")
plt.plot(x, 1 - np.tanh(kappa * x), 'k-.', linewidth=2, label="sigmoid")
plt.legend(prop={'size': 16})
plt.grid()
plt.xlabel('x', **axis_font)
plt.ylabel('$\phi(x)$', **axis_font)
```

(This problem is due to notes of Larry Wasserman.)

## 2.

Consider the least-squares problem with  $n$   $p$ -dimensional covariates,  $\{\mathbf{x}_i, y_i\}_{i=1}^n \subset \mathbb{R}^p \times \mathbb{R}$ . We would like to fit the following linear model,  $\hat{y}(\mathbf{x}) = \mathbf{x}^T \boldsymbol{\beta}$ . Also, suppose that there are coefficients  $C_+ \subset \{1, \dots, p\}$  such that for all  $j \in C_+$  we require that  $\beta_j \geq 0$ ,  $j \in C_+$ , and another set  $C_- \subset \{1, \dots, p\}$ , such that  $\beta_j \leq 0$ ,  $j \in C_-$  (assume that  $C_+$  and  $C_-$  are non-overlapping). Suppose that  $\mathbf{X}^T \mathbf{X}$  is invertible.

Such examples occur in insurance applications: the cost of a given insurance policy is based on a model for the amount of money a customer will cost the company, and each covariate is a variable specific to the customer (such as gender, age, credit history, etc.). It looks bad for the company if the insurance policy is more expensive for a customer that has an older account with the company than a newer account, when everything else is held fixed. Let  $x_{i,j} = 1\{\text{customer } i \text{ has had a policy for more than 2 years}\}$ , then  $\beta_j \geq 0$  is necessary for this property to hold.

(a) Write the constrained optimization for the empirical risk minimization with the constraints.

**Answer:**

$$\min_{\beta} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 \quad (5a)$$

$$\text{s.t. } \beta_j \geq 0, j \in C_+ \quad (5b)$$

$$\beta_j \leq 0, j \in C_- \quad (5c)$$

(b) Derive the dual for the optimization as a function of dual parameters.

**Answer:** The Lagrangian is given by

$$L(\beta, \lambda) = \frac{1}{2} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \beta^T \lambda \quad (6)$$

where  $\lambda_j \geq 0$  for  $j \in C_-$ ,  $\lambda_j \leq 0$  for  $j \in C_+$  (note for  $j \in C_- \cap C_+$  we simply have  $\beta_j = 0$  thus  $\beta_j$  and the  $j$ -th column of  $\mathbf{X}$  can be removed from the problem) and  $\lambda_j = 0$  for  $j \notin C_+ \cup C_-$ . Set derivative to 0, we have

$$\frac{\partial L}{\partial \beta} = -\mathbf{X}^T(\mathbf{y} - \mathbf{X}\beta) + \lambda = 0 \quad (7)$$

therefore

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1}(\mathbf{X}^T \mathbf{y} - \lambda) \quad (8)$$

substitute back into (7), the dual function is

$$l(\lambda) = \frac{1}{2} \mathbf{y}^T \mathbf{y} - \frac{1}{2} (\mathbf{X}^T \mathbf{y} - \lambda)^T (\mathbf{X}^T \mathbf{X})^{-1} (\mathbf{X}^T \mathbf{y} - \lambda) \quad (9)$$

(c) Write the KKT conditions and remark on the implication of the complementary slackness condition.

**Answer:** The KKT conditions are

$$-\mathbf{X}^T(\mathbf{y} - \mathbf{X}\beta) + \lambda = 0; \text{ (stationarity)} \quad (10a)$$

$$\beta_j \geq 0, j \in C_+; \beta_j \leq 0, j \in C_-; \text{ (primal feasibility)} \quad (10b)$$

$$\lambda_j \leq 0, j \in C_+; \lambda_j \geq 0, j \in C_-; \lambda_j = 0, j \notin C_+ \cap C_-; \text{ (dual feasibility)} \quad (10c)$$

$$\lambda_j \beta_j = 0. \text{ (complementary slackness)} \quad (10d)$$

The complementary slackness condition implies that, if  $\lambda_j \neq 0$  we must have  $\beta_j = 0$ , and if  $\beta_j \neq 0$  we must have  $\lambda_j = 0$ .

### 3.

Look at the dataset which can be found here: <https://archive.ics.uci.edu/ml/datasets/Blood+Transfusion+Service+Center>

(a) You will predict the final column in the dataset, which is an indicator if the person has made a blood donation. Form three different kernel functions and  $n \times n$  kernel matrices of  $K_{i,j} = k(\mathbf{x}_i, \mathbf{x}_j)$  that you think might be appropriate. You may use 2 different built in kernels, but must define one yourself.

**Answer:** We will use two-built in kernels from scikit-learn, namely the linear kernel and the radial basis function kernel. We will also use the Epanechnikov kernel function

$$K(x, x') = D \left( \frac{|x - x'|}{\lambda} \right) \quad (11)$$

Table 1: Comparison of different classifiers.

Classifier	Optimal parameters	Training score	Test score
SVC+Linear	$C = 10^4$	0.7754	0.7326
SVC+RBF	$\gamma = 10^{-2}, C = 10^2$	0.7840	0.7433
SVC+Epanechnikov	$l = 10^2, C = 10^3$	0.7743	0.7326
KNN+Linear	$k = 2$	0.7786	0.7380
KNN+RBF	$k = 2, \gamma = 10^{-4}$	0.7797	0.7380
KNN+Epanechnikov	$k = 2, l = 10$	0.7797	0.7380

where

$$D(t) = \begin{cases} \frac{3}{4}(1 - t^2) & |t| < 1, \\ 0 & \text{otherwise.} \end{cases} \quad (12)$$

(b) Apply kernel SVMs and  $k$ -nearest neighbors (where the distance is  $d(i, j) = k(\mathbf{x}_i, \mathbf{x}_i) + k(\mathbf{x}_j, \mathbf{x}_j) - 2k(\mathbf{x}_i, \mathbf{x}_j)$ ) with each of the 3 kernels.

**Answer:** See next section.

(c) Tune any parameters based on what you have learned about validation, and compare these methods with test errors (there are 6 different methods to compare, kNN and SVM with each kernel).

**Answer:** We split the dataset into a 25% test set and a 75% training+validation set. The parameters of the rbf kernel and the Epanechnikov kernel are tuned based on a Stratified ShuffleSplit cross validation, where the test\_size is set to 1/3 of the training+validation dataset (25% of the original dataset) with 5 iterations.

The range of tuning parameters are set as follows. For the support vector classifier, we take  $\log C = \text{range}(-4, 4)$ . For the  $k$ -nearest neighbor classifier we take  $k = \text{range}(1, 15)$ . For the rbf kernel, we take  $\log \gamma = \text{range}(-4, 4)$ . For the Epanechnikov kernel, we take bandwidth  $\log l = \text{range}(-3, 3)$ .

The tuning and testing results of the 6 tuned classifiers shown in Table. 1. These results suggest that all the classifiers perform almost equally badly, as a trivial classifier always output 0 would achieve almost the same score over the given dataset.

```
import pandas as pd
import numpy as np
import timeit
import sys

from sklearn import preprocessing, neighbors
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.cross_validation import train_test_split,
    ↳ StratifiedShuffleSplit
from sklearn.grid_search import GridSearchCV
from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.pipeline import Pipeline

df = pd.read_csv('transfusion.data') # Import the datafile
```

```

X = preprocessing.scale(df.values[:, :4].astype("float64")) # The
    ↪ predictors, standardized
y = df.values[:, 4] # The responses
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25)
    ↪ # Set aside 25% of the data for test
print("{0} is {0}, {1} is {1}".format(1 - sum(y) / len(y), sum(y) / len(y)))

cv = StratifiedShuffleSplit(y_train, n_iter=5, test_size=1/3, random_state
    ↪ =42) # Generate the cross validation labels

C_range = np.logspace(-4, 4, 9) # Range of Parameter C for all 3 svcs
k_range = np.arange(1, 15) # Range of parameter n_neighbors for all 3
    ↪ knncls

gamma_range = np.logspace(-4, 4, 9) # Range of parameter gamma for rbf
    ↪ kernel
l_range = np.logspace(-3, 3, 7) # Range of parameter l for Epanechnikov
    ↪ kernel

classifiers = dict() # The set of classifiers
params_grid = dict() # The set of cross validation parameters for each
    ↪ classifier

# The svm classifier with linear kernel
classifiers["svc_linear"] = SVC(kernel='linear')
params_grid["svc_linear"] = dict(C=C_range)

# The svm classifier with rbf kernel
classifiers["svc_rbf"] = SVC(kernel='rbf')
params_grid["svc_rbf"] = dict(gamma=gamma_range, C=C_range)

# The svm classifier with user defined Epanechnikov kernel
def epa(X, Y, l=1.0):
    """Epanechnikov kernel with bandwidth parameter
    return a n_row_X-by-n_row_Y matrix
    """
    n_sample_X = len(X)
    n_sample_Y = len(Y)
    t = np.empty([n_sample_X, n_sample_Y], dtype="float64")
    for i in range(n_sample_X):
        for j in range(n_sample_Y):
            t[i, j] = np.linalg.norm(X[i] - Y[j]) / l

    return 3 / 4 * (1 - t ** 2) * (abs(t) < 1)

class EpaKernel(BaseEstimator, TransformerMixin):
    def __init__(self, l=1.0):
        super(EpaKernel, self).__init__()
        self.l = l

```

```

def transform(self, X):
    return epa(X, self.X_train_, l=self.l)

def fit(self, X, y=None, **fit_params):
    self.X_train_ = X
    return self

classifiers["svc_epa"] = Pipeline([('epa', EpKernel()), ('svm', SVC()),])
params_grid["svc_epa"] = dict([('epa__l', l_range), ('svm__kernel', [
    ↪ 'precomputed']), ('svm__C', C_range),])

# The knn classifier with linear kernel
classifiers["knnc_linear"] = neighbors.KNeighborsClassifier()
params_grid["knnc_linear"] = dict(n_neighbors=k_range)

# The knn classifier with rbf kernel
def kernel_rbf(x, y, gamma):
    return np.exp(-gamma * np.linalg.norm(x-y) ** 2)

def dist_rbf(x, y, gamma):
    return kernel_rbf(x, x, gamma) + kernel_rbf(y, y, gamma) - 2 *
    ↪ kernel_rbf(x, y, gamma)

classifiers["knnc_rbf"] = neighbors.KNeighborsClassifier(metric=dist_rbf)
params_grid["knnc_rbf"] = dict(metric_params=[{'gamma': gamma} for gamma
    ↪ in gamma_range], n_neighbors=k_range)

# The knn classifier with user defined Epanechnikov kernel
def kernel_epa(x, y, l):
    t = np.linalg.norm(x - y) / l
    return 3 / 4 * (1 - t ** 2) if abs(t) < 1 else 0.0

def dist_epa(x, y, l):
    return kernel_epa(x, x, l) + kernel_epa(y, y, l) - 2 * kernel_epa(x, y
    ↪ , l)

classifiers["knnc_epa"] = neighbors.KNeighborsClassifier(metric=dist_epa)
params_grid["knnc_epa"] = dict(metric_params=[{'l': l} for l in l_range],
    ↪ n_neighbors=k_range)

for classifier in classifiers.keys():
    start_time = timeit.default_timer()

    print("Tuning_{}_...".format(classifier))

    grid = GridSearchCV(classifiers[classifier], param_grid=params_grid[
    ↪ classifier], cv=cv, verbose=0, n_jobs=4)
    grid.fit(X_train, y_train)

```



```
print("The best parameters are {0} with a score of {1}".format(grid
    ↪ .best_params_, grid.best_score_))
print("The score on the test set is {0}".format(grid.score(X_test,
    ↪ y_test)))
print("Elapsed time is {0}".format(timeit.default_timer() -
    ↪ start_time))
sys.stdout.flush()
```